

- Dependability and Fault Tolerance
- System Modelling
- Replication by Software
- Reliable Group Communication
- Distributed Agreement
- Consistent Database Replication



- Software based replication
 - Introduction
 - Consistency criteria
 - Active replication
 - Passive replication




- Replication is **the** solution to achieve fault tolerance.
- Software based replication is an **economically appealing** technology to provide generic fault-tolerant services.
- Replication is usually expected to be transparent to the user. However, transparency impacts on performance and it's not always desirable.
- Services are replicated to increase their **dependability** and often their **performance**. However, these goals are often conflicting.
- Replica **consistency** is a key issue. Consistency enables reliability, availability and performance trade-offs.




Introduction

Non-replicated server



Replicated server



- How consistent shall the replicas be?
- What aspects shall the replication protocol favor?
- What fault model shall be considered?



- How consistent shall the replicas be?
- How to express replica consistency?
- The requirement of **strictly consistent replicas** could be stated as:


Any read operation returns the value corresponding to the most **recent** written value

- Relies on absolute global time. Cannot be achieved in a distributed system
- Consistency criteria are usually stated as a set of predicates over system's executions




- A replicated service is said to be **Linearizable** if for every execution E there is some **interleaved sequence** S of all operations of E such that the following properties are satisfied:
- S satisfies the semantics of the non replicated service
- The **realtime** order of non-concurrent operations in E is preserved in S .


- Consider a FIFO queue:



- A replicated service is said to be **Sequentially Consistent** if for every execution E there is some **interleaved sequence** S of all operations of E such that the following properties are satisfied:
 - S satisfies the semantics of the non replicated service
 - The **local** order of operations in E is preserved in S .
- Consider a FIFO queue:



- In a replicated system (even) a non sequentially consistent execution can be easily produced even in the absence of faults:




- To satisfy Linearizability two conditions are necessary and sufficient:

Order: Given invocations $op(arg)$ and $op'(arg')$ on replicated server x , any two replicas of x that handle both operations do so in the same order.

Atomicity: Given invocation $op(arg)$ by client p on replicated server x , if one replica of x handles the operation, then every non faulty replica of x also handles $op(arg)$.




Active Replication



1. Request goes to all the replicas

2. Each replica processes the request, updates its own state, and returns the response to the client




3. Client waits until it receives the first response (or a majority of identical responses)

- Active replication requires request execution to be **deterministic**.
- The failure of a replica tends to be transparent for the clients.
- The failure and reintegration of the replicas has however to be dealt by the replication protocol.
- The active replication protocol is based on a communication primitive **available to clients** that ensures the required order and atomicity properties.
- This primitive is called **total-order multicast** or **atomic multicast**




Passive Replication



1. Request goes to a distinguished replica - the primary

2. The primary processes the request, updates its own state, and sends a state update message to all other replicas.




3. Each replica updates its own state, and sends an acknowledgment to the primary.


4. The primary returns the response to the client.



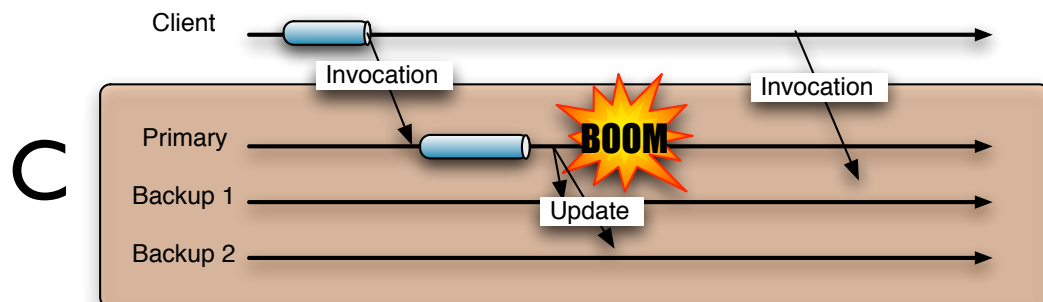
Passive Replication



Transparent to the client



Client reissues invocation



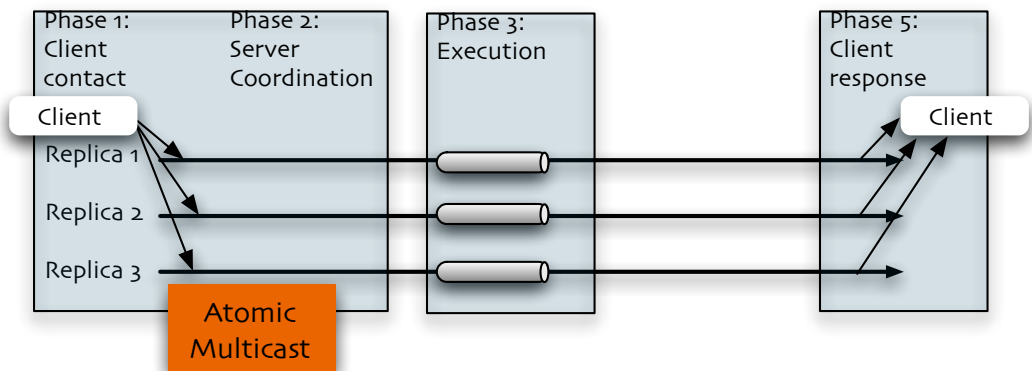

Which backups received updates?
How does the new primary process the reissued invocation?



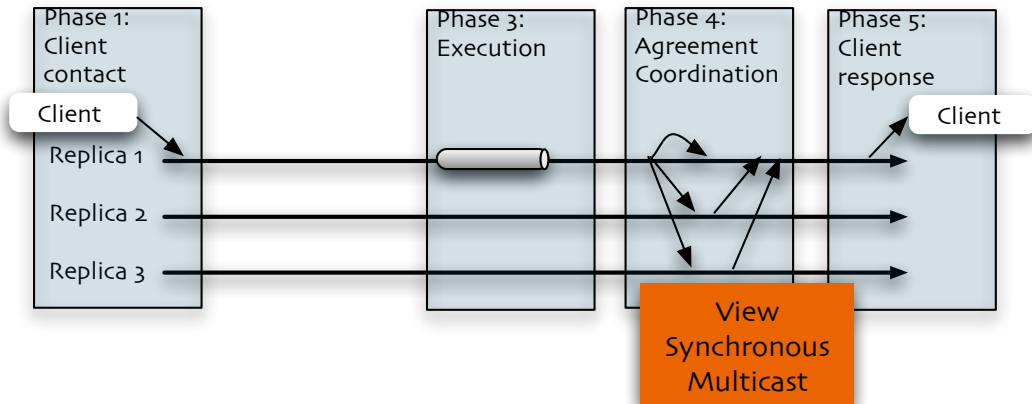
- Passive replication is asymmetric: Only one replica actually processes the requests. Execution does not need to be **deterministic**.
- The failure of the primary may not be transparent to the clients.
- The choice of the primary, as well as the failure and reintegration of the replicas has to be dealt by the replication protocol.
- The passive replication protocol is based on a **group membership** service and on the **view-synchronous multicast** communication primitive.



Functional model of replication



Active Replication



Passive Replication

